

122COM: Performance and Scalability

David Croft

Coventry University

david.croft@coventry.ac.uk

2015

Overview

1 Profiling

2 $O()$ notation

- Simple algorithms
- Good algorithms
- Bad algorithms

3 Recap

For any large piece of code you should:

- Write clear, easily understood code. Focus on getting the behaviour right, not on performance.
- Test the performance.
 - It may be fine.
- Profile your code to get the baseline performance.
 - So that you know if you are making things better or worse.
- Focus your efforts on the code that is consuming all the time.
 - E.g. small pieces of code that get called multiple times.

Statistical

- Regularly checks the system state.
- Will slow down running speed.
 - But equally throughout the code.

Instrumental

- bb

Flat Profiler

■ a

Call Profiler

■ b

Profiler types

O() notation

Profiling is very useful in determining the actual performance of your code.

- Not so good at measuring how code will scale.
- Algorithmic complexity.
- Certain algorithms are known to be better than other algorithms.

O() notation

Used to describe complexity in terms of time and/or space.

- Commonly encountered examples...
 - $O(1)$, $O(\log n)$, $O(n)$, $O(n \log n)$, $O(n^2)$,
 $O(2^n)$ and $O(n!)$
- n refers to the number of values.
 - E.g. n values to be sorted.
 - E.g. n values to be searched.

$O(n)$

Linear time.

- n is directly related to time/space required.
- E.g. linear search.

$O()$ notation describes the worst case scenario.

- Usually otherwise stated.

| | | | | | | | | |
|------------|---|---|---|---|---|---|---|---|
| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
| array[n] = | A | B | C | D | E | F | G | H |

$O(1)$

Constant time.

- n doesn't matter.
- Always takes same time/space.
- E.g. getting first item in an array.

$O(n^2)$

Square of the elements.



- A lot of sorting algorithms are $O(n^2)$.
- Nested `for` loops

```
for i in range(n):  
    for j in range(n):  
        pass
```
- $O(n^3)$, $O(n^4)$, $O(n^m)$ etc. are all possible.

$O(\log n)$

Logarithmic time.



- Bit more complicated.
- Binary search.



$O(\log n)$

Logarithmic time.


- Bit more complicated.
- Binary search.



$O(\log n)$


Logarithmic time.

- Bit more complicated.
- Binary search.



$$n = 3$$


$$O(\log n) = 1.58 \Rightarrow 1$$



$O(\log n)$


Logarithmic time.

- Bit more complicated.
- Binary search.



$$n = 3$$

$$O(\log n) = 1.58 \Rightarrow 1$$




$$n = 7$$

$O(\log n)$


Logarithmic time.

- Bit more complicated.
- Binary search.



$$n = 3$$

$$O(\log n) = 1.58 \Rightarrow 1$$




$$n = 7$$

$$O() = 2.81 \Rightarrow 2$$

$O(\log n)$ cont.

$O(\log n)$ complexity increases very slowly.

- $O()$ for a hundred items is only 6.
- $O()$ for a trillion item is only 39!



$$n = 15$$

$$O() = 3.91 \Rightarrow 3$$

$O(n \log n)$

More logarithmic time.

- Looks more difficult than it is.
- $O(n \log n)$ means, do $O(\log n)$ n times.
- E.g. binary search for n items.
 - Binary search is $O(\log n)$.
 - Doing n binary searches.
 - So $O(n \log n)$.

$$O(2^n)$$

Exponential time.

- Very, very bad.
- Each additional value doubles the time/space.
- Doesn't scale.

Factorial time.

$$O(n!)$$

Factorial time.

$$O(n!)$$

$O(n!)$

Factorial time.


- The worst.
- Every possible combination of n items.
- Brute force travelling salesman is $O(n!)$.
- Totally impractical even for small values of n .

Relative performance

Different $O()$ == wildly different complexity.

| | | n | | |
|-------|---------------|-----|---------|-----------------------|
| | | 2 | 10 | 100 |
| Best | $O(1)$ | 1 | 1 | 1 |
| | $O(\log n)$ | 1 | 3 | 6 |
| | $O(n)$ | 2 | 10 | 100 |
| | ↑ | | | |
| | ↓ | | | |
| Worst | $O(n \log n)$ | 2 | 33 | 664 |
| | $O(n^2)$ | 4 | 100 | 10000 |
| | $O(2^n)$ | 4 | 1024 | $1.27 \cdot 10^{30}$ |
| | $O(n!)$ | 2 | 3628800 | $9.33 \cdot 10^{157}$ |

Comparison



Scalability

Scalability isn't efficiency.

- A good $O(n^2)$ implementation can be better than a bad $O(n)$.
 - For a while.
- Eventually, as n increases, $O(n)$ will always outperform $O(n^2)$ etc.

Recap

Profiling determines the actual performance of your code.

■ .

$O()$ notation is how your code should scale.

Recap

Profiling determines the actual performance of your code.

- .

$O()$ notation is how your code should scale.

- Theoretically.

Recap

Profiling determines the actual performance of your code.

- .

$O()$ notation is how your code should scale.

- Theoretically.
- Lots of real world issues can mess it up.

The End